MDM EXPRESS · INTERNAL TECHNICAL SPECIFICATION

# Onboarding Agent Module

The queue, cadence, capacity model and workspace that take a new merchant from first registration to an activated account — built as a ticket type inside the Omnichannel Support Platform, not as a system of its own.

**HOW TO READ THIS DOCUMENT**

Parts 0 and 1 are for every reader and define scope, vocabulary and the one architectural decision everything else depends on. Parts 2 to 13 are the engineering specification. Every reader should read **Part 0.4 (Guiding principles)** and **Part 12 (Configuration surface)**, because together they define what Operations can change without a deployment. **Part 15 is the acceptance criteria** — the testable definition of done.

**THIS DOCUMENT IS SUBORDINATE**

The *Omnichannel Support Platform — Technical Specification v1.3* is the parent document and the platform this module is built on. Where the two disagree, **the platform specification wins and the discrepancy is a defect in this document**. Nothing here introduces a new datastore, a new routing engine or a second agent application; every entity named below is either an existing platform entity or a configuration row inside one.

| DOCUMENT | STATUS | AUDIENCE |
|---|---|---|
| **Technical Specification v1.0** | **Ready for build — 3 decisions outstanding** | **Engineering, Product, Operations, QA** |

# Contents

**Part 0 — About This Document**

Purpose, audience, the delivery package, glossary, guiding principles, out of scope.

**Part 1 — System Overview**

What the module is, its relationship to the support platform, system context, the core mental model.

**Part 2 — Data Model**

The `onboarding_task` ticket type, its field set, the status machine, the handoff record.

**Part 3 — Capture, Identity & Migration**

Where leads come from, how a registration becomes a ticket, duplicate handling, decommissioning Zoho.

**Part 4 — Assignment & Capacity**

The availability formula, the cost of a lead, the selection algorithm, the holding queue, overrides.

**Part 5 — The Cadence Engine**

The two cadences, reset-on-reply, cutoffs, win-back, and the interaction with the template budget.

**Part 6 — The Agent Workspace**

Screen inventory, the docked conversation, macros, templates and the session window.

**Part 7 — Walkthrough & Resource Hub**

The 20-minute call, booking, the declined-call path, and the self-serve alternative.

**Part 8 — Activation & Handoff**

The qualification gate, the accept action, ownership transfer to the Account Manager.

**Part 9 — Roles & Permissions**

The Onboarding Agent role, the administrator, cross-team visibility, field-level rules.

**Part 10 — The Time Model**

Minute values, the handle-time formula, and the limits on what it may be used for.

**Part 11 — Metrics & Reporting**

Why onboarding is excluded from support metrics, the five OA formulas, dashboards, comparability.

## Part 12 — Configuration Surface

What Operations can change without code, and what requires engineering.

## Part 13 — The Admin Workspace

Operations, performance, conversation monitoring, change management, audit.

## Part 14 — Delivery Plan & Open Decisions

Phases, dependencies on the platform build, and the decisions that block them.

## Part 15 — Acceptance Criteria

The testable definition of done for every feature area.

## Appendices

A. Permission matrix · B. Rule catalogue · C. Metric formulas · D. Configuration parameters · E. Decision log.

# Part 0 — About This Document

## 0.1 Purpose

This document specifies the Onboarding Agent module: the work surface, queue, cadence engine, capacity model and reporting that carry a newly registered merchant from their first message to an activated account handed to an Account Manager.

It is written to be implemented from directly. Every entity, rule, permission, formula and acceptance criterion named here is intended to be built as described, or challenged explicitly before build starts.

It replaces the current Zoho CRM onboarding pipeline. That pipeline is not migrated feature by feature: its *responsibilities* are absorbed into the support platform, and the lead-tracking logic it accumulated is expressed as configuration inside the routing engine specified in the parent document. **After cutover there is no CRM in the onboarding path.**

## 0.2 Audience and how to read

| Reader | Read | Why |
|---|---|---|
| **Management / Operations** | Parts 0, 1, 4, 5, 12, 14 | Scope, how leads are distributed, what the system does on its own, what you can change yourself, and the decisions still waiting on you. |
| **Product** | Parts 0–9, 11–14 | Full behavioural specification excluding infrastructure detail. |
| **Engineering** | All parts | Parts 2, 4, 5, 9, 10, 12 carry the binding technical contracts. |
| **QA** | Parts 2, 4, 5, 9, 11, 15, Appendices B–D | The state machine, the capacity arithmetic, rule precedence and metric definitions are where defects concentrate. Part 15 is the test basis. |

### 0.2.1 The delivery package

Four artefacts travel together. Their standing differs and confusing them causes rework.

| Artefact | Standing | Use it for |
|---|---|---|
| **Omnichannel Support Platform** `MDM-Omnichannel-Support-Platform-Spec.pdf` v1.3 | **Normative — and superior** | The platform this module runs inside. Contacts, conversations, tickets, the routing engine, channels, permissions, the audit log and the agent workspace are all defined there. This document adds a ticket type and a set of views; it does not redefine any of that. |

| Artefact | Standing | Use it for |
|---|---|---|
| **This specification** `MDM-Onboarding-Agent-Module-Spec.pdf` v1.0 | **Normative** | Every onboarding-specific entity, rule, permission, minute value, metric and acceptance criterion. Disagreements inside the onboarding domain are resolved here. |
| **OA console prototype** `mdm-oa-console.html` | Behavioural reference | Screen inventory, layout intent, the agent and admin role split, the docked conversation, the cadence tracker, and the draft → simulate → publish flow. Open it and switch roles before writing code. Where the prototype and this text disagree, this text wins and the prototype is the defect. |
| **Onboarding Agent role & playbook** `MDM_Onboarding_Agent_Role_v1_1.pdf` | Business process | What the human does and why — call scripts, objection handling, qualification judgement, escalation, compensation. It defines the job. This document defines the software that supports it. |

> **ON COLOURS AND TYPOGRAPHY**
>
> The prototype is rendered in the MDM brand palette because it is intended to be shown to the onboarding team. Where colour carries meaning — session-window state, lead status, capacity pressure, ownership — *the meaning is normative and the hue is not*.

## 0.3 Glossary

Each term below is used in this document with exactly one meaning. Terms already defined in the platform specification (Contact, Conversation, Ticket, Team, Agent, Role, Macro, Template, Rule) keep their definitions from there and are not repeated.

| Term | Definition in this document |
|---|---|
| **Lead** | Not an entity. A convenience word for the triple of one Contact, one Conversation and one open `onboarding_task` Ticket. There is no lead table. |
| **Onboarding task** | The ticket type that carries onboarding work. The unit of assignment, of measurement and of completion. |
| **Cadence** | A configured sequence of scheduled outreach attempts with a fixed cutoff. Two exist: the no-answer cadence and the went-quiet cadence. |
| **Attempt** | One outbound call on the no-answer cadence, whether answered or not. Counted in `attempt_no`. |
| **Follow-up** | One outbound message on the went-quiet cadence. Counted in `followup_no`. Always a template, therefore always billable. |

| Term | Definition in this document |
|---|---|
| **Walkthrough** | The scheduled 20-minute call in which the merchant is shown the platform and the pricing calculator with their own numbers. |
| **Handoff record** | The eight fields the Account Manager inherits at activation. Its completeness is a measured KPI. |
| **Holding queue** | Onboarding tickets with no assignee because no agent had capacity, or because they arrived outside business hours. A state the system passes through, never rests in. |
| **Assignable minutes** | An agent's shift length minus the reserved triage and close blocks. The denominator of availability. |
| **Free minutes** | Assignable minutes minus today's committed work. The number the assignment engine sorts on. |
| **Slot** | Capacity for one more lead. Bounded by both free minutes and the open-lead ceiling; the smaller bound governs. |
| **Session window** | The 24-hour period after a merchant's last inbound message during which free-text replies are permitted. Outside it, only approved templates may be sent. |
| **Handle time** | Estimated agent minutes spent on a lead, *derived* from the ticket's own record using configured minute values. Never measured by a timer and never self-reported. |

## 0.4 Guiding principles

These are binding. Where a later section conflicts with a principle, the principle wins and the section is a defect.

**PRINCIPLE 1 — THIS IS A TICKET TYPE, NOT AN APPLICATION**

The parent specification prohibits building a second interface for a second audience. Onboarding is one of its four named audiences. This module therefore adds one configured ticket type, one team, one role, a set of saved views and a typed rendering of that ticket — and nothing else. No new datastore, no second inbox, no parallel routing engine, no duplicate status vocabulary.

**PRINCIPLE 2 — THE PLATFORM ASSIGNS; AGENTS NEVER CHOOSE**

Every lead has an owner from the moment it exists. There is no claim action anywhere in the agent interface. Cherry-picking is not a behaviour to be discouraged by policy; it is a capability that is absent by construction. Where no agent can take a lead, the system holds it with a stated reason and assigns it automatically when that reason clears.

**PRINCIPLE 3 — ONBOARDING IS MEASURED BY COMPLETION, NOT BY RESPONSE TIME**

Onboarding work is proactive and sequenced. It carries no SLA policy, is excluded from CSAT collection, and contributes to no support volume, response or resolution aggregate. Its ticket type carries `counts_toward: none` for exactly this reason, and that value is not to be changed without amending this document.

**PRINCIPLE 4 — THE CONVERSATION IS NEVER DUPLICATED**

The onboarding workspace shows work state and the platform conversation side by side, as one instance of one component. There is no second composer, no second template picker, no second delivery-state display and no second place an agent might type. Closing an onboarding ticket never closes the conversation; the relationship outlives the task.

**PRINCIPLE 5 — EVERY TOUCH HAS A PRICE, AND THE SYSTEM KNOWS IT**

A call costs minutes. A template costs money. A follow-up costs both. The cadence, the capacity model and the reporting all read from the same configured price list, so a change to the process shows its cost before it is published rather than at the end of the month.

**PRINCIPLE 6 — EVERYTHING THE BUSINESS WILL WANT DIFFERENT IN SIX MONTHS IS CONFIGURATION**

Attempt counts, gaps, cutoffs, minute values, capacity ceilings, template content, budgets, statuses, permitted transitions, permissions and every user-visible label are configuration. A change of onboarding process must never require a deployment.

## 0.5 Out of scope

- **Post-activation account management.** The module ends at a complete handoff. The 30-day playbook, the merchant's book and the AM's own metrics belong to the Account Management system.
- **Order, pricing and settlement logic.** Read and displayed, never written. The order domain owns all of it.
- **Marketing acquisition.** Ad targeting, landing pages and campaign messaging are upstream. This module reports on source quality; it does not act on it.
- **The Resource Hub content management system.** Article authoring, versioning and publishing are the platform knowledge base (parent spec, Part 12). This document specifies only how onboarding consumes and shares it.
- **Recruitment, compensation and coaching process.** Defined in the role playbook. The module supplies the numbers; it does not implement the bonus calculation.

# Part 1 — System Overview

## 1.1 What the module is

A single work surface for the Onboarding team, in which every new merchant is a sequenced piece of work with an owner, a next action, a deadline and a complete audit trail — and in which the conversation with that merchant is present, live, in the same screen.

Three properties distinguish it from the CRM pipeline it replaces:

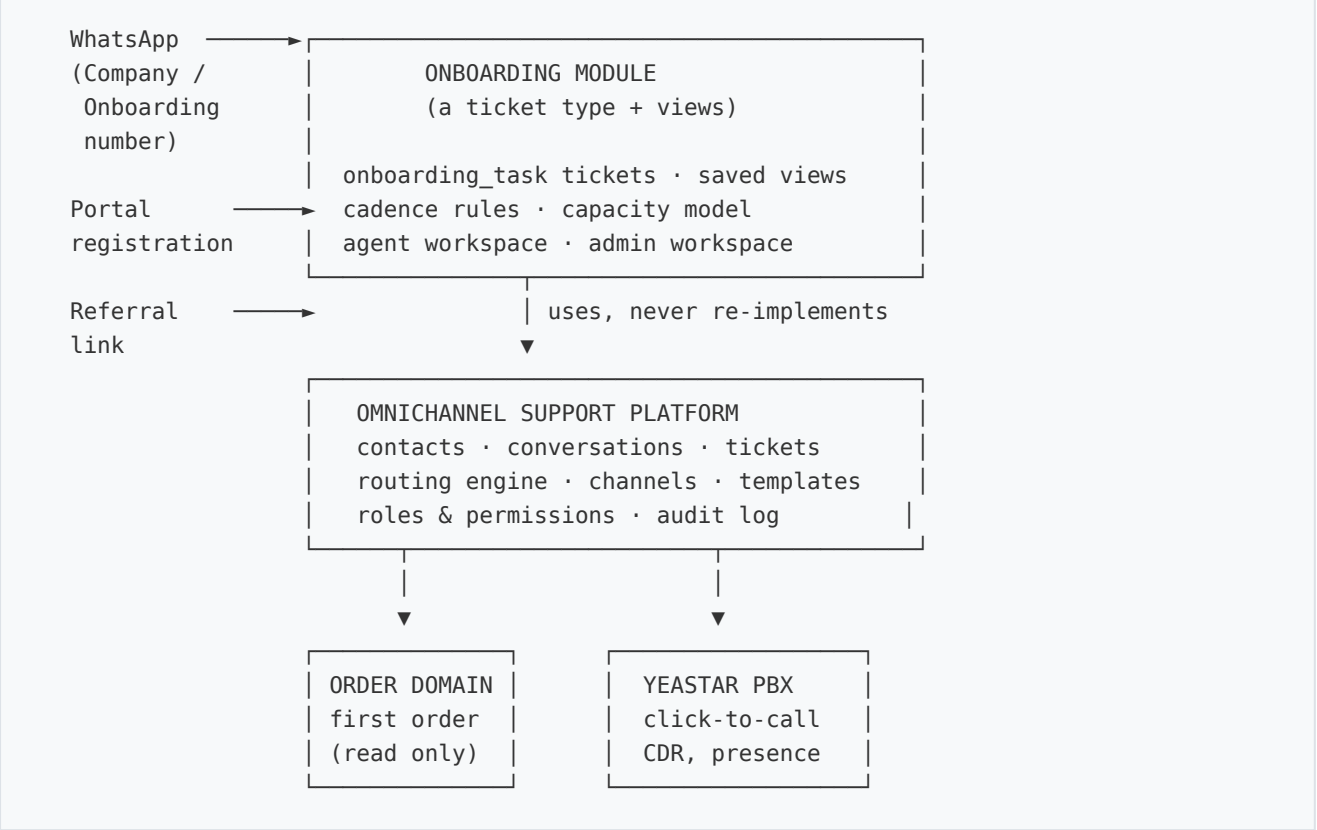
- **Distribution is computed, not chosen.** Leads are assigned to the agent with the most free working minutes on that market. Nobody picks, and nobody is handed work they cannot absorb.
- **The process runs itself.** Reminders, escalations, follow-up sends and cutoffs are rules in the platform's routing engine. An agent who forgets a lead does not lose it.
- **Work and conversation are one screen.** The agent never navigates between "the CRM record" and "the messaging tool", because there is only one of each and they are docked together.

## 1.2 The audience this module serves

The parent specification names four audiences. This module serves exactly one of them.

| Audience | Reaches us via | Owning team | Defining characteristic |
|---|---|---|---|
| **New merchant** | Company number (voice + WhatsApp), portal registration | **Onboarding** | Work is proactive and sequenced, not reactive. Measured by completion, not response time. Relationship is temporary — it ends at activation. |
| End customer | WhatsApp, voice | Customer Service | Out of scope here. |
| Merchant (activated) | Portal widget, WhatsApp | Merchant Support | Out of scope. An activated merchant who messages the onboarding number is forwarded to their AM the same day and never advised by the OA. |
| Internal team | In-system only | Warehouse, Finance, IT | Out of scope. |

## 1.3 System context




The module is a module. Its boundary with the support platform is an internal one — it calls the platform's ticket, conversation, routing and permission services and stores nothing of its own outside the configured field set in Part 2.

## 1.4 The core mental model

Three objects, and one relationship between them that must not be broken.

| Object | In onboarding it is | Lifespan | Owned by |
|---|---|---|---|
| **Contact** | The merchant — permanently, before and after activation | Permanent | Support platform |
| **Conversation** | The thread with them, across WhatsApp, portal and voice | Permanent | Support platform |
| **Onboarding task** | The work of getting them started | Opened once, closed once | This module |

**WORKED ILLUSTRATION**

Fatima registers through a referral link on Sunday. One **contact** is created and one **conversation** opens on her WhatsApp number. One **onboarding task** is created and assigned to Salma. Salma calls three times over three days without an answer; the task carries `attempt_no = 3` and the conversation carries three call events and one template send. On Wednesday Fatima replies. The cadence resets, the walkthrough is booked, and on Friday the task is closed as activated — **the conversation stays open and moves with Fatima to her Account Manager**. Six weeks later she complains about a late parcel: same contact, same conversation, a new ticket of a different type owned by a different team.

## 1.5 The journey

Five stages, carried in `ticket.field[stage]`. Stage is the position in the process; status (Part 2.4) is the state of the relationship. Both are needed, and conflating them is the most common modelling error in this domain.

| Stage | Entered when | The agent's next action |
|---|---|---|
| `first_contact` | The ticket is created | Call within the first-contact SLA. The goal of this call is a booked walkthrough, not a sale. |
| `walkthrough` | A walkthrough is booked, or the merchant is sent to the Resource Hub instead | Deliver the 20-minute call, or run the went-quiet cadence on the self-serve path. |
| `documents` | The merchant is ready but a required document is missing | Chase the specific missing item. Nothing else is blocking. |
| `activation` | All five qualification criteria are met | Perform the accept action and write the handoff record. |
| `done` | Activation complete and handoff recorded | None. The ticket closes; the AM takes the contact. |

## 1.6 Why this is not a CRM

The instinct to solve onboarding with a sales CRM should be resisted, for three concrete reasons.

1. **The conversation is the same conversation.** A merchant messages one number before and after activation. In a separate CRM, the pre-activation thread is stranded in a system the Account Manager cannot see, and the merchant is asked to repeat themselves on day one of the relationship.
2. **Field visibility becomes a code risk instead of a configuration.** An onboarding agent must not see merchant cost, margin or settlement. In the platform this is one field-level permission bound to `contact_type`, set once. In two systems it is a rule living in two codebases that will eventually diverge.

3. **Capacity cannot be computed across two systems.** An agent's real load is calls, walkthroughs and messages. A CRM that cannot see the conversation cannot see the work, and will distribute leads by count while the actual constraint is minutes.

# Part 2 — Data Model

The recurring theme of the parent specification holds here without exception: **types are data, not code**. This module introduces no new tables. It introduces one row in `ticket_type`, its field set, its status set and its permitted transitions — all administrator-editable.

## 2.1 What this module adds

| Added | Where it lives | Notes |
|---|---|---|
| Ticket type `onboarding_task` | `ticket_type` table | One configuration row. Already present in the platform's launch seed catalogue. |
| Its field set (2.3) | `ticket_type.fields` | Custom fields on the existing ticket entity. No schema migration. |
| Its status set and transitions (2.4, 2.5) | `status` configuration | Statuses are rows; transitions are rows. |
| Team `Onboarding` | `team` | Customer-facing. One queue. |
| Role `Onboarding Agent` | `role` | Grants in Part 9 and Appendix A. |
| Ten routing rules (Appendix B) | `rule` | Composed from the platform's existing condition and action catalogue. |
| Five WhatsApp templates (7.5) | WABA template registry | Utility category. Meta approval required before use. |
| Six macros (6.5) | `macro`, team scope | Internal. Insert text and apply changes; never send. |
| Saved views (6.3) | `view`, system and team scope | Stored filters. A view is a query, not a folder. |

Nothing else. In particular there is no `lead` entity, no `pipeline`, no `stage_history` table (the event log already carries it) and no onboarding-specific message store.

## 2.2 Ticket type

| Attribute | Value |
|---|---|
| `key` | `onboarding_task` |
| `visibility` | `external` — the contact may be messaged on this ticket |
| `default_team` | `Onboarding` |
| `sla_policy` | `null` — no SLA. Timing is governed by cadence due dates, not by a response clock. |

| Attribute | Value |
|---|---|
| `counts_toward` | `none` |
| `survey_eligible` | `false` |
| `one_open_per_contact` | `true` — a contact may hold at most one open onboarding task at a time |

> **COUNTS_TOWARD MUST STAY AT NONE**
>
> This single value is what keeps onboarding work out of the customer-service volume, SLA compliance and CSAT figures. An onboarding task has no SLA to breach and a merchant who has not started trading has nothing to rate. Setting it to anything else silently corrupts every support metric in the platform, and the corruption is not visible on any screen until someone asks why response times moved.

## 2.3 Field set

| Key | Type | Written by | Purpose |
|---|---|---|---|
| `stage` | select | Agent, rules | Position in the journey (1.5). Values: `first_contact · walkthrough · documents · activation · done`. |
| `due_date` | datetime | Rules | The next action's deadline. Drives every reminder and the overdue count. |
| `attempt_no` | integer | Agent action | No-answer cadence counter. Incremented when a call outcome is logged, never by a rule. |
| `next_attempt_at` | datetime | Rules | When the next call reminder fires. |
| `followup_no` | integer | Rules | Went-quiet cadence counter. Incremented by rule `r5` when a follow-up template is sent. |
| `next_followup_at` | datetime | Rules | When the next follow-up template is due. |
| `cadence_reset` | boolean | Rules | Guards the reset-on-reply rule so it fires once per ticket, not repeatedly (5.4). |
| `hold_reason` | text | Rules | Why an unassigned ticket is in the holding queue. Never null while unassigned (4.6). |
| `walkthrough_at` | datetime | Agent | Scheduled start of the walkthrough call. |
| `walkthrough_mode` | select | Agent | `phone · meet · whatsapp_video`. |
| `qualification` | boolean × 5 | Agent | The five activation criteria (8.1). All five true is the precondition for the accept action. |
| `handoff_record` | composite | Agent | Eight sub-fields (2.6). Required to reach `activated`. |
| `source_utm` | text | Capture | Channel plus UTM or affiliate tag, carried from registration. Immutable after creation. |

| Key | Type | Written by | Purpose |
|---|---|---|---|
| `first_order_at` | datetime | **Order domain** | Written by a domain event, never by this module. Feeds the shared first-order KPI. |

## 2.4 Status set

Statuses use the vocabulary the onboarding team already uses. Each maps to one of the platform's status categories, and the category — not the label — drives reporting and clock behaviour.

| Status | Category | Meaning |
|---|---|---|
| `new` — New | `open` | Created and assigned. No contact attempted yet. |
| `failed_to_contact` — Failed to Contact | `open` | One or more attempts made, no reply. Carries `attempt_no` as a visible counter on the status chip. |
| `prequalified` — Pre-qualified | `open` | Reached and interested, but not committed. |
| `qualified` — Qualified | `open` | Reached, real intent, ready to be walked through and activated. |
| `activated` — Activated | `solved` | Account live, handoff record complete, AM assigned. |
| `cold_lead` — Cold Lead | `closed` | A cadence reached its cutoff, or the merchant declined. Excluded from the activation-rate denominator only if also junk. |
| `junk_lead` — Junk Lead | `closed` | Not a real merchant: wrong number, duplicate, test, unreachable by construction. Excluded from every metric. |

> **ONE STATUS, ONE COUNTER — NOT FIVE STATUSES**
>
> Repeat contact attempts are **not** separate statuses. A second failed call does not move the ticket to "second failed attempt"; it increments `attempt_no` while the status stays `failed_to_contact`. The interface renders this as a numeric badge on the status chip. This keeps the status set at seven rows instead of eleven, makes the cutoff a single numeric condition, and means changing the number of attempts is a configuration edit rather than a status redesign.

## 2.5 State machine

Illegal transitions are rejected by the service, not merely hidden in the interface.

| From | May move to | Required on transition |
|---|---|---|
| New | Failed to Contact · Pre-qualified · Qualified · Junk Lead | — |

| From | May move to | Required on transition |
|---|---|---|
| Failed to Contact | Pre-qualified · Qualified · Cold Lead · Junk Lead | To Cold: a reason, or the cadence cutoff having been reached |
| Pre-qualified | Qualified · Failed to Contact · Cold Lead · Junk Lead | To Cold: a written reason |
| Qualified | **Activated** · Failed to Contact · Cold Lead | To Activated: all five qualification criteria true *and* a complete handoff record |
| Activated | — (terminal for this ticket type) | — |
| Cold Lead | Pre-qualified (reopen) | A written reason, recorded against the actor |
| Junk Lead | Pre-qualified (reopen) | A written reason, recorded against the actor |

Note the absence of a direct `New → Activated` path. The fast-track case — an experienced seller who wants to start immediately — is still routed through Qualified, because the qualification criteria and the handoff record are what make an activation defensible, and speed is not a reason to skip the record.

## 2.6 The handoff record

Eight fields. Activation without a complete record counts against the clean-handoff rate (11.2) and is visible to the Account Manager as an incomplete inheritance.

| Field | Content |
|---|---|
| Account Manager | The named AM who becomes the contact's persistent owner at activation. |
| Seller profile | Who they are, city, experience, what they sold on before. |
| Intent | Products, target market or markets, expected monthly volume. |
| Services discussed | Which services they care about, and why. |
| Objections & answers | What they pushed back on and what convinced them. |
| Commitments made | **Every promise, verbatim.** The AM must honour all of it. |
| Risk notes | Anything that needs watching. |
| Source | Channel plus UTM or affiliate tag, copied from `source_utm`. |

## 2.7 What this module never writes

- **Order tables.** `first_order_at` arrives as a domain event. The order card is read-only, and cost, margin and settlement fields are absent from the API response for this role, not merely hidden.

- **Seller account state beyond the accept action.** Activation calls the existing platform accept operation; it does not set account flags directly.
- **Contact ownership after activation.** The transfer of `contact.owner_agent_id` to the AM is performed by the platform as part of the accept action.
- **Its own audit trail.** Every state change, assignment, field edit, permission use and outbound message goes to the platform's single immutable event log.

# Part 3 — Capture, Identity & Migration

## 3.1 Sources

All onboarding traffic arrives on the Company / Onboarding number or through the platform's registration event. Numbers are allocated per function, never per person, so a number survives a staff change and a lead can be reassigned without the merchant noticing.

| Source | Arrives as | Identity strength |
|---|---|---|
| WhatsApp — Company / Onboarding line | Inbound message | Phone number, E.164 — medium |
| Portal registration | `contact.registered` domain event | Portal session — strong |
| Voice — Company / Onboarding line | Inbound call event (Yeastar) | Caller ID — weak until confirmed on the call |
| Referral / affiliate link | Registration event with affiliate tag preserved | Portal session — strong |
| Forwarded social lead | Manual ticket creation by the OA, tagged with source | Whatever the social channel supplied — weak |

## 3.2 The capture rule

This is rule `r1` (Appendix B) and it is the mechanism that replaces the CRM's list.

```
WHEN   message.received on the Onboarding number
   OR   contact.registered

IF     contact.contact_type = merchant
 AND   contact has no open ticket of type onboarding_task

THEN   create_ticket(type: onboarding_task)
       set_team(Onboarding)
       assign_agent(strategy: most_free_minutes, skill: contact.country)
       set_field(stage, first_contact)
       set_field(due_date, now + first_contact_sla)
       set_field(source_utm, event.source + event.utm)
       notify(assignee)
```

Two consequences worth stating plainly. First, a merchant who messages before registering and registers afterwards produces **one** ticket, because the second event finds an open one. Second, an activated merchant who messages the onboarding number produces **no**

onboarding ticket — their `contact_type` has moved on, and the message is routed to their Account Manager instead.

## 3.3 Duplicate registrations and identity

Identity resolution is the platform's, not this module's. Three onboarding-specific consequences:

- **Same phone, second registration.** Resolves to the existing contact. If an open onboarding task exists, no new ticket is created and the existing one is annotated. If the contact is already activated, escalate — a second account for an existing seller is checked with that seller's AM before anything is activated.
- **Merge conflicts are queued for human review**, never resolved silently. A merged contact keeps both conversations' history.
- **Suspected fake registration is not an OA decision.** The agent marks Junk Lead only for the unambiguous cases — number not in service, obvious test data, duplicate of their own record. Suspected fraud is escalated to the Commercial Director with evidence, and the ticket is neither activated nor rejected in the meantime.

## 3.4 Arrivals outside business hours

Registrations do not wait for an agent to notice them, but they are also not assigned to somebody who is not working. A ticket created outside the Onboarding team's business hours is created, assigned `null`, and given `hold_reason = "outside business hours"`. Rule `r8` assigns it at the first sweep after the shift opens. The first-contact clock starts at ticket creation and is measured in business hours only.

## 3.5 Decommissioning Zoho

The CRM is not run in parallel. There is a single cutover, and it is reversible up to the point the WhatsApp number moves.

| Item | Approach |
|---|---|
| Contacts | Full import, de-duplicated by phone number through the platform's identity rules. Merge conflicts queued for human review. |
| Open leads | Created as `onboarding_task` tickets with their original creation timestamps. **Cadence counters do not carry over** — every migrated lead starts at `attempt_no = 0` with a due date at cutover, because the old counters were maintained by hand and cannot be trusted. |
| Closed / lost leads | Imported as closed tickets for searchability. Excluded from all metrics before the cutover date. |
| Message history | Imported into conversations, read-only, tagged `legacy`. This is what stops agents asking merchants to repeat themselves in week one. |

| Item | Approach |
|---|---|
| Status mapping | Zoho statuses map to the seven in 2.4 one-to-one. Any Zoho status with no mapping is imported as `prequalified` with a note, never dropped. |
| The onboarding WhatsApp number | Moves in a single cutover. It cannot serve both systems at once. |
| Zoho access | Read-only for two weeks after cutover, then exported and archived per the retention policy. |

**BEFORE THE CUTOVER DATE, NOT AFTER**

The five WhatsApp templates in 7.5 require Meta approval, which takes time that is outside our control. Submit them at the start of Phase 1. A cutover that arrives before template approval leaves the team unable to contact any merchant whose 24-hour window has closed — which, on the first morning, is most of them.

# Part 4 — Assignment & Capacity

## 4.1 The rule that has no exceptions

Leads are assigned by the platform. There is no claim action, no unassigned pool an agent can browse, and no "available leads" view. An agent opens their queue and finds work that is already theirs.

This is a capability decision, not a policy one. Cherry-picking is impossible because the button does not exist, which also removes the collision case entirely — two agents cannot claim the same lead if neither can claim anything.

## 4.2 Availability — the formula

Availability answers one question: *how much more work can this agent absorb today?* It is measured in minutes, not in head-count, because onboarding load is calls and walkthroughs, and two agents holding five leads each can have entirely different days.

```
assignable   = shift_minutes − reserved_minutes

committed    = walkthroughs_today  × (walkthrough_minutes + walkthrough_buffer)
             + first_contacts_due  × (first_contact_minutes + wrap_minutes)
             + followups_due       ×  followup_minutes

free         = max(0, assignable − committed)

lead_cost    = first_contact_minutes + wrap_minutes + (followup_minutes × 2)

slots        = 0                                              if agent is on a call
             = min( floor(free ÷ lead_cost),
                    max_open_leads − open_leads )               otherwise

can_take     = slots > 0
```

### 4.2.1 What each term counts

| Term | Counts |
|---|---|
| `walkthroughs_today` | Open tickets with `stage = walkthrough` and `walkthrough_at` falling today. |
| `first_contacts_due` | Open tickets with `status = new`. Every one of these is an unmade first call. |
| `followups_due` | Open tickets not in the two sets above, whose `due_date` is today or overdue. |

| Term | Counts |
|---|---|
| `open_leads` | All open tickets assigned to the agent, regardless of whether anything is due today. |

### 4.2.2 Worked example

Default parameters (Appendix D): shift 480, reserved 60, walkthrough 20 + 5, first contact 6 + 3, follow-up 4, ceiling 12.

```
Salma — Libya queue, 5 open leads

assignable = 480 − 60                        = 420 min
committed  = 1 walkthrough  × 25             =  25
           + 2 first contacts × 9            =  18
           + 2 follow-ups     × 4            =   8
                                               ———
                                                51 min
free       = 420 − 51                        = 369 min
lead_cost  = 6 + 3 + (4 × 2)                 =  17 min
slots      = min( floor(369 ÷ 17), 12 − 5 )
           = min( 21, 7 )                    =   7   ← bound by the ceiling
```

> **TWO CONSTRAINTS, AND THE INTERFACE NAMES WHICH ONE BINDS**
>
> At MDM's current volume — roughly 70 registrations a month across three agents — the binding constraint is almost always the open-lead ceiling, not the minute budget. That is a finding, not a defect: it means overdue items are a cadence-discipline problem rather than a capacity problem, and that hiring would not fix them. The Operations screen must state which constraint is binding for each agent, because a manager who reads "six hours free" and "cannot take more" without the explanation will conclude the system is broken.

### 4.2.3 Why a lead costs more than its first call

`lead_cost` is deliberately not the first-contact call alone. Accepting a lead commits the agent to the cadence that follows it — on the current defaults, an expected two further touches before the outcome is known. Pricing only the first call would let the engine fill an agent's day with new leads whose real cost lands tomorrow, which is precisely the failure mode round-robin produces.

## 4.3 Selection

```
function pick_agent(ticket):
    pool = agents in team Onboarding
             where skill.country contains ticket.contact.country
    if pool is empty:  pool = all agents in team Onboarding

    eligible = [a for a in pool if availability(a).can_take]
    if eligible is empty:  return NULL          → holding queue (4.4)

    sort eligible by  free_minutes  DESC,
                      open_leads    ASC          (tie-break)
    return eligible[0]
```

Country is a *skill*, not a hard partition. If no agent carries the market, the whole team is eligible — a lead is never held because of a language preference when somebody is free.

## 4.4 The holding queue

When `pick_agent` returns null the ticket is created with no assignee and a mandatory `hold_reason`. It is a state the system passes through, never rests in.

| Hold reason | Cleared when |
|---|---|
| `outside business hours` | The first sweep after the shift opens. |
| `no agent with free capacity` | Any agent's free minutes or open-lead count changes such that `can_take` becomes true. |
| `all agents on calls` | A call ends. |

Rule `r8` sweeps the holding queue on a configured interval and on the `agent.capacity_changed` event. If a ticket has been held longer than the configured escalation threshold, the shift lead is notified — the holding queue is allowed to exist, but it is not allowed to be quiet.

## 4.5 Overrides

| Action | Who | Behaviour |
|---|---|---|
| Assign over capacity | Administrator, shift lead | Assigns a held ticket to the agent with the most free minutes, ignoring `can_take`. Agents on calls are still excluded. Audited. |
| Reassign | Administrator, shift lead | Moves ownership. The merchant sees nothing; the conversation is unbroken. Audited with before and after. |
| Bulk reassign | Administrator, shift lead | Up to the role's bulk ceiling. Per-ticket permissions are evaluated individually and partial failure is reported honestly. |

| Action | Who | Behaviour |
|---|---|---|
| **Request transfer** | Onboarding Agent | The agent's only assignment-adjacent action. Sends a request with a written reason to the shift lead. The agent cannot move a ticket themselves, including onto themselves. |
| Take over | Administrator, shift lead | Reassignment initiated from the conversation monitor (13.4). Requires a reason. The previous owner is notified. |

## 4.6 What availability is not

- **Not a productivity score.** It is a forward-looking budget of today's remaining time. An agent with low free minutes is busy, not good; an agent with high free minutes is available, not idle.
- **Not a per-channel concurrency count.** The platform's per-channel capacity model still applies to reactive conversation routing. Onboarding assignment uses minutes because onboarding work is scheduled rather than queued.
- **Not visible to the agent as a target.** The agent's own screen shows their committed and free minutes so they can plan the day. It carries no goal, no comparison to peers and no colour that implies underperformance.

# Part 5 — The Cadence Engine

Two cadences, each a counter, a gap and a cutoff. Both are configuration. Together they are what replaces a human remembering to follow up.

## 5.1 The two cadences

| | No-answer cadence | Went-quiet cadence |
|---|---|---|
| **Applies when** | The merchant has never replied | The merchant replied once, then stopped |
| **Touch** | A phone call, with a WhatsApp nudge where the window allows | A WhatsApp template |
| **Counter** | `attempt_no` | `followup_no` |
| **Default limit** | 5 attempts | 3 follow-ups |
| **Default gap** | 1 day | 3 days |
| **Status while running** | `failed_to_contact` | Unchanged — usually `prequalified` |
| **Cutoff** | `cold_lead` | `cold_lead` |
| **Cost** | Agent minutes; templates only when the window is shut | **Always billable** — every follow-up lands outside the window by construction |

## 5.2 The no-answer cadence

The counter is incremented by the *agent* logging a call outcome. The rules create the reminder and enforce the cutoff; they never fabricate an attempt that nobody made.

```
on agent logs call outcome = no_answer:
    attempt_no       += 1                    (capped at no_answer_max)
    status            = failed_to_contact
    if attempt_no >= no_answer_max:
        status        = cold_lead
        due_date      = null
        log reason: "N daily attempts, no answer — cadence cutoff"
    else:
        next_attempt_at = now + no_answer_gap
        due_date        = next_attempt_at
```

The status **stays** `failed_to_contact` for the whole run. Only a reply changes it, and only by hand — an agent judges whether a merchant who finally answers is Qualified, Pre-qualified or Junk. The system never infers intent from the fact of a reply.

## 5.3 The went-quiet cadence

Here the system does send. Rule `r5` fires on status age, sends an approved template, and increments the counter.

```
WHEN   ticket.status_age_exceeded (awaiting reply > quiet_gap)

IF     followup_no < quiet_max
 AND   contact.status != opted_out
 AND   template_budget.remaining > 0

THEN   send_message(whatsapp, template: oa_nudge_v2)
       set_field(followup_no, +1)
       set_field(next_followup_at, now + quiet_gap)
       set_field(due_date, next_followup_at)
```

**THE BUDGET CONDITION IS LOAD-BEARING**

`template_budget.remaining > 0` is not defensive decoration. A scheduled rule that sends outside the session window with no budget check will exhaust the monthly ceiling silently, and the first sign will be a Meta invoice. Keep the condition on the rule, and keep the per-contact cap (Appendix D) as a second line of defence.

## 5.4 Reset on reply — once

An inbound message from a merchant on a running cadence resets it. The reset happens **once per ticket**, guarded by `cadence_reset`.

```
WHEN   message.received (inbound)
IF     ticket.type = onboarding_task
 AND   ticket.cadence_reset = false
THEN   set_field(attempt_no, 0)
       set_field(followup_no, 0)
       set_field(cadence_reset, true)
       set_field(due_date, now)
       notify(assignee)
```

Without the guard, a merchant who replies "ok" every few days and never engages resets the clock indefinitely and never reaches a cutoff. With it, the second disengagement runs to the cutoff on schedule.

## 5.5 Cutoff, win-back and archive

| Event | Default | Behaviour |
|---|---|---|
| Cutoff | Attempt 5 / follow-up 3 | Status → Cold Lead, `due_date` cleared, reason logged, assignee notified. The ticket leaves the working queue. |
| Stalled flag | Day 7 | Any ticket still open and not yet activated at day 7 appears on the admin funnel view and requires a written reason. It does not change status. |
| Win-back | Day 30 | One template to cold leads, then archive. **Disabled by default** — see the note below. |
| Archive | Day 30 + 7 | Ticket closed. Contact and conversation are untouched and remain searchable. |

> **WHY WIN-BACK SHIPS DISABLED**
>
> The cohort evidence available at the time of writing shows activation after day 14 is close to zero. A day-30 win-back therefore spends template budget on the least responsive population. Rule `r9` is specified, built and shipped *off*; enabling it is a one-click configuration change and should be made deliberately, with the cohort table open.

## 5.6 Precedence

Where more than one rule matches, the platform's rule precedence applies. Within onboarding the intended order is:

1. `r7` reset-on-reply — an inbound message always takes precedence over a scheduled send.
2. `r4` / `r6` cutoffs — a ticket at its limit is closed before another touch is scheduled.
3. `r5` follow-up send.
4. `r3` attempt reminder.
5. `r2` first-contact escalation.

The practical case this ordering exists for: a merchant replies at 08:59 on the morning their third follow-up is due. They must receive the agent, not the template.

## 5.7 Cadence parameters

All eight are administrator-editable and take effect at the next rule evaluation. None of them is retroactive: a ticket already closed is not reopened by raising a limit.

| Parameter | Default | Effect of changing it |
|---|---|---|
| `first_contact_sla` | 24 h | Moves the first-contact due date and the escalation threshold in `r2`. |
| `no_answer_max` | 5 | Lowering it moves currently-open tickets at or above the new limit to Cold Lead at the next sweep. |

| Parameter | Default | Effect of changing it |
|---|---|---|
| `no_answer_gap` | 1 day | Spacing of call reminders. |
| `quiet_max` | 3 | Directly multiplies template spend on the quiet cohort. |
| `quiet_gap` | 3 days | Spacing of follow-up sends. |
| `stalled_day` | 7 | When a written reason becomes required. |
| `winback_day` | 30 | Only meaningful while `r9` is enabled. |
| `template_budget_month` | See D | Hard stop for `r5` and every other automated send. |

# Part 6 — The Agent Workspace

## 6.1 Screen inventory

| Screen | Purpose |
|---|---|
| **My Day** | Landing screen. Own KPIs, cadence due now, today's walkthroughs, capacity, and leads newly assigned by the platform. |
| **Onboarding Queue** | All onboarding tickets, filtered through saved views. Read for the whole team, editable only for own. |
| **Lead workspace** | Work state and the live conversation, side by side (6.4). The screen the agent spends the day in. |
| **Reminders** | Cadence items grouped overdue / today / tomorrow / this week. |
| **Walkthroughs** | Scheduled calls, the self-serve cohort, and the 20-minute agenda. |
| **Resource Hub** | Read and share. Authoring is a knowledge-base permission the agent does not hold. |
| **Playbook & Macros** | Scripts, objection frames, the macro list, approved templates, and a read-only statement of how the queue runs. |
| **My Performance** | Own metrics only, plus own quality reviews with a right of reply. |

There is no settings screen in the agent workspace. Configuration is administrator territory (Part 12); what an agent needs to *know* about the rules is presented read-only inside the Playbook, because an agent who does not know what the system will do on its own will duplicate it by hand.

## 6.2 Visibility and editability

An Onboarding Agent can open every onboarding ticket and read it in full. They can act only on their own. This is deliberate: the team works one shared number, and an agent who cannot see a colleague's thread cannot cover for them, answer a merchant who calls back, or learn from a good walkthrough.

| On a ticket that is not theirs | Behaviour |
|---|---|
| Reading everything | Permitted — ticket, conversation, history, handoff record. |
| Adding an internal note | Permitted. This is how coverage and context-sharing happen. |
| Replying to the merchant | Blocked. The composer is replaced by the note composer with the reason stated. |
| Status change, macro, booking, activation | Blocked. Every control is *disabled and labelled with why*, never hidden. |
| Reassigning it to themselves | Blocked. "Request transfer" is offered instead. |

> **DISABLED, WITH THE REASON PRINTED**
>
> A disabled control must always state its own reason, in the place the agent is looking. "Activate account — 2 qualification criteria still unmet" and "Run a macro — lead not assigned to you" are correct. A greyed button with a tooltip, or a control that silently disappears, is not: both produce a support ticket instead of an understanding.

## 6.3 Saved views

A view is a stored filter with a name, a sort order and a column set. A ticket appears in every view whose conditions it satisfies and is never moved between them. Counts are live.

| Level | Managed by | Onboarding views at launch |
|---|---|---|
| System | Fixed | Assigned to me · Due today · Held in queue · All leads |
| Team | Manager | Attempting contact · Went quiet · Walkthroughs · Activated · Cold & junk |
| Personal | The agent | Unlimited. The one configuration surface an agent owns. |

A shared view shows each viewer only what their own role permits, so row counts legitimately differ between two people looking at the same view.

### 6.3.1 Queue columns

The queue's default column set is specified because one column in it is not obvious and is the single most useful signal on the screen.

| Column | Content |
|---|---|
| Lead | Name, Arabic name, ticket and contact reference. |
| Channel | Origin channel and market. |
| **Session window** | **Time remaining on the 24-hour window, or "template only".** Whether the next touch is free or costs money, visible before the ticket is opened. |
| Status | Status chip with the attempt counter where applicable. |
| Stage | Position in the journey. |
| Next action | Due date, rendered as overdue / today / dated. |
| First order | Present only after activation. From the order domain. |
| Owner | Assignee, with own tickets visually distinguished. |

## 6.4 The lead workspace — the docked conversation

This is the integration point between this module and the platform, and the reason the module does not need a messaging surface of its own.

The screen is two panes. The left pane is work state: status and stage, cadence tracker, qualification checklist, order card, contact record, handoff record and ticket history. The right pane is **the platform conversation component**, embedded — the same component the Customer Service and Merchant Support workspaces render, with the same composer, the same session-window banner, the same template picker, the same delivery states and the same collision indicators.

| Requirement | Detail |
|---|---|
| One instance, not a copy | The conversation pane is the platform component with a layout variant. It is not a reimplementation, and no onboarding-specific message rendering, composer or send path exists. |
| Session window is shown before writing | Remaining window time is displayed in the pane header at all times, not discovered when a send fails. |
| Free text is blocked server-side outside the window | The composer replaces itself with the approved-template picker. The block is enforced in the service; removing the control from the interface is not sufficient. |
| Deep link out | "Open in OmniDesk" opens the same conversation in the full agent workspace, for the cases where an agent wants multi-conversation triage. |
| Deep link in | A conversation opened in OmniDesk that has an open onboarding task shows a link to that ticket in its context rail. The navigation is symmetric. |
| Collision detection | Inherited from the platform: viewer indicators, typing indicators, stale-reply warning, and prevention of two agents replying blind. |
| Draft auto-save | Per ticket, surviving navigation and reload. Inherited. |

## 6.5 Macros

A macro applies a set of ticket changes and inserts text into the composer. It never sends. Each macro states what it will do before it runs.

| Macro | Applies |
|---|---|
| Log call — no answer | `attempt_no` +1 · status → Failed to Contact · reminder at `+no_answer_gap` · internal note |
| Log call — answered, strong lead | `attempt_no` +1 · status → Qualified · stage → walkthrough · internal note |
| Log call — answered, soft lead | `attempt_no` +1 · status → Pre-qualified · internal note |
| Book the walkthrough | stage → walkthrough · sets `walkthrough_at` and mode · inserts confirmation text · due date at call time |
| Send Hub instead of a call | stage → walkthrough (self-serve) · sends `oa_hub_share_v1` · starts the went-quiet cadence |

| Macro | Applies |
|---|---|
| Mark junk — not a real seller | status → Junk Lead · clears due date · internal note |

Macros are scoped globally, per team or personally. The six above are **team** macros maintained by the manager; an agent may create personal macros for themselves and nothing wider. Usage counts are kept, so an unused macro can be retired and a macro run two hundred times a month can be considered for a rule.

## 6.6 Templates and the session window

A macro and a template are different objects and the interface must not blur them. A macro is free and internal; a template costs money and needs Meta approval. The distinction is the single most common source of agent confusion in systems of this kind, so:

- Templates are presented only where they are required — inside a closed session window — and always with the word *billable* attached to the send.
- A template send shows the resolved message with variables filled from the contact record before it goes.
- Sending with an unapproved or rejected template fails loudly at configuration time, not silently at send time.
- All outbound goes under the single MDM Express sender identity, from the Company / Onboarding number, regardless of which merchant it concerns.

## 6.7 What the agent cannot do

- Pick, claim or reassign a lead, including onto themselves.
- Reject a merchant. Suspected fraud escalates; it is not an OA decision.
- Author or edit a template, a team macro, a rule, a knowledge-base article or any budget.
- Send a bulk outbound message. Bulk internal changes and bulk outbound are separate permissions and only the first is granted.
- See merchant cost, platform fee or settlement fields — absent from the API response, with the interface stating they are withheld rather than showing blanks.
- Improvise pricing, or hint that an exception is possible.

# Part 7 — Walkthrough & Resource Hub

## 7.1 The walkthrough

A scheduled **20-minute** call, plus a 5-minute buffer for notes and the handoff record. Twenty-five minutes is the number the capacity model prices (Part 4) and the reporting estimates (Part 10); changing the call length changes both, immediately.

| Min | Block | Content |
|---|---|---|
| 0–3 | Their story | Recap discovery in one line and let them correct you. What does success look like in 90 days? |
| 3–7 | Services mapped | Only the services that serve their goal. Name two, not seven. |
| **7–13** | **Pricing, live** | Open the calculator with *their* product and *their* volumes. Full cost per delivered order, no surprises. The longest block, because it is the trust moment. |
| 13–16 | Process end to end | Sourcing request to COD payout at headline level. Detail lives in the Resource Hub. |
| 16–18 | After activation | "Your Account Manager calls within 24 hours and stays with you daily for your first month." |
| 18–20 | Close | Ready → activate on the call. Hesitant → handle the objection, then one concrete next step with a date. |

> **WHY 20 AND NOT 45**
>
> The earlier 45-minute structure spent half its time on process detail that is better read than heard, and it capped an agent at roughly eight walkthroughs a day. At 20 + 5 the same agent day holds sixteen. Pricing keeps the largest share because it is the block that converts; the process detail moves to the Resource Hub rather than being cut, so nothing is lost — it is relocated to where a merchant can re-read it.

## 7.2 Booking

- Two named slots are offered, never an open question. The interface presents slot buttons, not a free-text field, because "when suits you?" is the most common reason a booking never happens.
- Booking sets `walkthrough_at`, `walkthrough_mode` and the ticket's due date, and moves stage to `walkthrough`.
- The confirmation goes out on the merchant's own channel, with the Resource Hub link attached, so they arrive with questions rather than from cold.
- If the session window is closed at booking time, the confirmation is sent as `oa_slot_proposal_v3` and is billable.

- A booked walkthrough consumes `walkthrough_minutes + buffer` of the agent's availability for that day from the moment it is booked.

## 7.3 The declined-call path

A merchant who does not want a call is not a lost merchant. Sending the Resource Hub is a first-class outcome, not a fallback:

```
stage        → walkthrough (self-serve)
send          oa_hub_share_v1
followup_no   = 0
due_date      = now + quiet_gap        → the went-quiet cadence begins
```

The qualification criterion "walkthrough delivered" may be satisfied by the Hub for a demonstrably experienced seller (8.1), so this path can still end in activation.

## 7.4 The Resource Hub

Everything the walkthrough covers, in writing, in Arabic and English. It is the platform knowledge base with an onboarding audience filter; this module specifies only the consumption side.

| Requirement | Detail |
|---|---|
| Chaptered recording | The recorded walkthrough is chaptered against the six blocks in 7.1, so an agent can point at a single chapter — "watch 7 to 13, that's the pricing" — instead of sending 20 minutes. |
| Section set at launch | How MDM works end to end · Pricing and the delivery-fee calculator · Sourcing without capital · COD, wallet and payouts · Why the confirmation call raises delivery rate · What happens after activation. |
| Sharing | One share action from the lead workspace, which sends `oa_hub_share_v1` and records the send on the ticket. |
| Ownership | Authored and maintained by the Senior OA under the knowledge-base permission. Agents have read and share only, and report outdated content rather than editing it. |
| Feedback loop | Repeated questions in onboarding conversations are the input to new sections. The driver report (11.3) is where they surface. |

## 7.5 Templates at launch

| Key | Category | Used for |
|---|---|---|
| `oa_intro_v2` | utility | Post-registration introduction after an unanswered first call. |
| `oa_slot_proposal_v3` | utility | Two proposed walkthrough slots, when the window is closed. |

| Key | Category | Used for |
|---|---|---|
| `oa_hub_share_v1` | utility | Resource Hub link, pricing page and recording. |
| `oa_nudge_v2` | utility | The went-quiet follow-up: "pricing, trust, or timing?" A multiple-choice question is answered more readily than an open one. |
| `oa_winback_d30` | utility | Day-30 win-back. Ships disabled (5.5). |

All five are utility category. Marketing-category templates are prohibited on the onboarding number: the category affects both Meta's quality rating and what the number may be used for, and mixing them puts the line at risk.

# Part 8 — Activation & Handoff

## 8.1 The qualification gate

The accept action is unavailable until all five criteria are true. They are stored as booleans on the ticket and set by the agent.

| Criterion | Satisfied when |
|---|---|
| Real, reachable person | At least one call answered; phone verified. |
| Real intent | They can name what they will sell and in which market, even roughly. |
| Walkthrough delivered | The call happened, *or* was explicitly waived by a seller who demonstrably knows the platform. |
| Documents complete | No pending missing-documents item. |
| Expectations set | They can answer "what happens after activation?" correctly. |

> **DO NOT ACTIVATE TO HIT A NUMBER**
>
> A fake or unready activation costs more than a lost lead: it burns Account Manager time, poisons the shared first-order KPI for both roles, and surfaces as a bad-activation flag anyway. The gate is enforced in the service, so an unqualified activation is not merely discouraged — it is rejected.

## 8.2 The accept action

```
preconditions:  all five qualification criteria = true
                handoff_record complete (8 fields)
                actor holds activate_seller on this ticket

on accept:      call platform accept operation (idempotent)
                assign Account Manager as contact.owner_agent_id
                status  → activated
                stage   → done
                due_date = null
                start AM Day 0
                notify(account_manager) with the handoff record
                write immutable audit entry
```

The AM welcome call must be scheduled within 24 hours, and the agent confirms this before the ticket closes. The conversation is **not** closed and does not move — it stays with the contact, now owned by the AM.

## 8.3 Fast-track

An experienced seller who wants to start immediately may be activated before the walkthrough is delivered. The walkthrough is then booked for *after* activation, noted in the handoff record, and the AM is informed. Speed for the motivated is the point of the fast-track; skipping the record is not permitted, and the qualification criterion is satisfied by waiver, not bypassed.

## 8.4 Bad-activation flags

An Account Manager may flag an activation as bad within 7 days. Flags are reviewed monthly with the Commercial Director. The handoff record is the agent's defence, which is the operational reason it is mandatory rather than encouraged — it should be written as though somebody will read it, because somebody will.

A flagged activation is not reversed. It is recorded, counted against the clean-handoff rate if the record was incomplete, and fed to the quality review queue.

# Part 9 — Roles & Permissions

Permissions are granted to roles, never to named individuals. There is no interface path and no data path for a personal exception. Every permission is enforced in the service layer: removing a control from the interface is not sufficient, and the equivalent API call by an unauthorised role is rejected.

## 9.1 The Onboarding Agent role

Built from the platform's *Agent* template with two deliberate widenings and several narrowings. It is not a new class of user.

| Difference from stock Agent | Reason |
|---|---|
| **Widened:** view all team tickets → `team` scope | The team works one shared number. An agent must be able to read a colleague's thread to cover a callback. |
| **Widened:** internal note on any readable ticket | Coverage and context-sharing. Notes are the safe write. |
| **Narrowed:** no reassign / transfer | Assignment is the platform's (Principle 2). "Request transfer" is offered instead. |
| **Narrowed:** no reject | Rejecting a merchant is a Commercial Director decision. |
| **Narrowed:** no bulk outbound | Separate permission from bulk internal changes; not granted. |
| **Narrowed:** reports scoped to `own` | Coaching tool, not a leaderboard. |
| **Field-level:** cost, fee, margin, settlement withheld | Bound to `contact_type`, set once, enforced in the API response. |

## 9.2 The Onboarding Administrator

Not a new role either — the platform *Manager* and *Administrator* grants, scoped to the Onboarding team. Two capabilities matter for this module and are called out because they are easy to get wrong:

- **Monitoring is not replying.** An administrator reads any onboarding conversation. To reply to a merchant they must first take the ticket over, with a reason. Replying as a third identity would break the merchant's single point of contact — the interface therefore offers "take over" and "internal note", and offers no composer.
- **Configuration is versioned and attributed.** Every change an administrator publishes carries author, timestamp and diff, and is reversible (13.5).

## 9.3 Temporary elevation

Cover for leave is time-bound: a covering role expires automatically at the configured end and the reversion is audited. There is no permanent widening for a two-week absence.

## 9.4 Permission summary

The full matrix is Appendix A. The two rows that most often get implemented wrongly:

| Permission | Onboarding Agent | Note |
|---|---|---|
| View all onboarding tickets | `team` — read only | Read is team-wide; *write* is own-only. These are two separate grants and conflating them is the defect to watch for. |
| Activate seller (accept) | Granted, gated by 8.1 | The gate is a service-layer precondition, not an interface state. |

# Part 10 — The Time Model

## 10.1 Derived, never measured

Handle time is estimated from what the ticket already records — messages, calls, walkthroughs, notes, activation — multiplied by configured minute values. There is no timer, no "start work" button and no agent self-reporting.

This is a deliberate trade of precision for trust. A stopwatch produces a number nobody believes and a behaviour everybody games; a derived estimate produces a number that is consistent across agents, cannot be inflated, and is good enough for the only two questions it needs to answer: how much capacity does the team have, and what does an activation actually cost.

## 10.2 Minute values

All administrator-editable. These same values drive the capacity model in Part 4, so changing one moves assignment and reporting together, which is the intended coupling.

| Touch | Default | Applied to |
|---|---|---|
| Walkthrough call | 20 min | A delivered walkthrough |
| Walkthrough buffer | 5 min | Notes and handoff record after the call |
| First contact call | 6 min | The first call to a new registrant |
| Call wrap-up | 3 min | Logging the outcome of any answered call |
| Unanswered call | 1.5 min | Dial, wait, log |
| Cadence follow-up | 4 min | One attempt or one template send, including the note |
| Writing a message | 2 min | A free-text outbound reply |
| Reading a message | 0.5 min | An inbound message |
| Sending a template | 0.5 min | Pre-written — it costs money, not time |
| Internal note | 1.5 min | Any internal note |
| Activation + handoff record | 12 min | Once, at activation |

## 10.3 The handle-time formula

```
messaging   = outbound_free_text × compose_minutes
            + inbound_messages   × read_minutes
            + template_sends     × template_minutes

calls       = Σ over logged calls:
                answered  → call_duration + wrap_minutes      (duration from Yeastar
CDR)
                no answer → no_answer_minutes

walkthrough = walkthrough_minutes + buffer_minutes            if delivered, else 0

other       = internal_notes × note_minutes
            + activation_minutes                              if activated

handle_time = messaging + calls + walkthrough + other
```

Where a call duration is not available from the CDR, the configured defaults stand in. Where the ticket records touches that predate the conversation import (migrated leads), only the counters are used and the estimate is marked as partial.

## 10.4 What it may and may not be used for

| | Use |
|---|---|
| **May** | Compute agent availability and drive assignment (Part 4). Report the cost of an activation and of a lost lead. Size the team against forecast volume. Show an agent where their own time went on a lead. |
| **May not** | Appear in an individual's performance score or compensation. Be compared between agents as a productivity measure. Be presented to a merchant. Be treated as an actual, auditable time record for any purpose — payroll, billing or dispute. |

**CALIBRATE BEFORE TRUSTING**

The default minute values are estimates made before the system had run. They should be calibrated against real Yeastar call durations and real message timestamps at the end of Phase 1, and the defaults in Appendix D updated. Until that calibration happens, handle time is directional and every screen presenting it says *estimated*.

# Part 11 — Metrics & Reporting

## 11.1 Onboarding is measured separately, on purpose

The parent specification defines onboarding negatively: `counts_toward: none`, no SLA policy, excluded from CSAT collection, and its metric set deferred. This part closes that gap.

The five metrics below are built only on the ticket's own record — its status machine, `stage`, `due_date` and the event log — and none of them enters a support aggregate.

## 11.2 The five onboarding metrics

| Metric | Weight | Formula and exclusions |
|---|---|---|
| **Activation rate** | 30% | `count(activated) ÷ count(reached a terminal state)`, over registrants aged ≤ 14 days. **Junk excluded from both numerator and denominator.** Cold leads remain in the denominator — losing a lead is a result, not an exclusion. |
| **Speed to first contact** | 20% | `count(first_outbound_at − created_at ≤ first_contact_sla) ÷ count(tickets)`, business hours only. Computed from the event log, *not* from the platform's first-response-time metric — that one starts at human assignment and pauses out of hours, which would flatter this number. |
| **Clean handoff rate** | 20% | `count(activations with all 8 handoff fields non-empty) ÷ count(activations)`. |
| **First order ≤ 7 days** | 15% | `count(first_order_at − activated_at ≤ 7d) ÷ count(activations)`. **Shared with the Account Manager.** Computed in MDM reporting, displayed here (11.4). |
| **Funnel hygiene** | 15% | `count(open tickets past due_date) ÷ count(open tickets)`, sampled at end of day. Lower is better; target zero. |

Report the median and the 90th percentile for any duration, never the mean. Publish the denominator beside every rate.

> **WHY ONE KPI IS SHARED**
>
> First order within 7 days belongs to the onboarding agent *and* the account manager. It makes activation quality permanently the OA's problem: they cannot win a quarter by activating unready merchants, and the AM cannot win one by letting good handoffs sit.

## 11.3 Dashboards

| Dashboard | Audience | Contents |
|---|---|---|
| My Performance | Agent, `own` scope | The five metrics with threshold / target / exceptional bands, own history, the onboarding-team median, and own quality reviews with a right of reply. No peer ranking. |
| Operations | Admin, live | Awaiting first contact and the oldest wait; due and overdue; open conversations and how many need a template; held queue with reasons; per-agent availability; template budget burn; rule activity in the last 24 h. |
| Performance | Admin, 90-day window | Weekly activation rate **with dated configuration changes marked on the axis**; agent scorecards; conversion by source and by market; cohort activation at day 7 / 14 / 30; loss reasons; handle time by outcome. |
| Drivers | Admin | Structured loss reasons from a controlled list. If "never answered" is 38% of losses, the fix is the first-contact channel mix, not a sixth attempt. |

> **DATE EVERY CONFIGURATION CHANGE ON THE TREND**
>
> The activation-rate chart must mark configuration changes on the time axis with their date and a one-line description. Without it, nobody can tell whether a rule change caused a movement, and the weekly improvement loop degenerates into opinion. This is the single most valuable feature on the Performance screen and it costs almost nothing to build, because the change log already exists.

## 11.4 Boundary with MDM reporting

Exactly one domain defines each metric. The onboarding module reports on *contacts, work and teams*. It does not compute order metrics.

| This module reports | MDM reporting computes |
|---|---|
| Registrants, activations, activation rate | First order, order volume, confirmation rate |
| Cadence outcomes, loss reasons, cold and junk rates | Delivery and return rates |
| Speed to first contact, clean handoff rate, funnel hygiene | Revenue and settlement |
| Handle time and capacity | Merchant lifetime value |

The shared first-order KPI is computed in MDM reporting and read here. Cross-domain reports — activations per marketing dinar, revenue per onboarding hour — are built in the MDM reporting layer, which reads this module's data and not the reverse.

## 11.5 Comparability

Metrics are comparable only within the same audience and dependency profile. An onboarding agent is compared to their own history and to the onboarding-team median — never to a Customer Service agent, whose work, contacts and clocks are different in kind. The reporting interface enforces this by defaulting to same-segment comparison.

# Part 12 — Configuration Surface

This part is the contract between Operations and Engineering for the onboarding domain. If something in the left column later requires a code change, the implementation has failed to meet this specification.

## 12.1 Configurable without code

| Area | What can be changed | Permission |
|---|---|---|
| Cadence parameters | Attempt count, follow-up count, both gaps, first-contact SLA, stalled day, win-back day | Manager |
| Capacity model | Shift length, reserved blocks, open-lead ceiling | Manager |
| Time model | Every minute value in 10.2 | Manager |
| Assignment strategy | Strategy, skill dimension, tie-break, holding-queue sweep interval and escalation threshold | `manage_automation` |
| Rules | All ten: triggers, conditions, actions, priority, enable / disable | `manage_automation` |
| Statuses | Add, rename, reorder, permitted transitions, fields required on transition | Administrator |
| Ticket type & fields | Field set, visibility, metric participation | Administrator |
| Qualification criteria | The five gate items — wording, count, which are mandatory | Administrator |
| Handoff record | Which of the eight fields are required to activate | Manager |
| Templates | Content, variables, per-rule budgets, monthly ceiling, per-contact cap, enable / disable | Manager |
| Macros | Team macros: content and actions. Personal macros: the agent | Manager / Agent |
| Team & skills | Membership, market skills, business hours, per-channel capacity | Manager |
| Roles & permissions | Grants, scopes, field-level rules, time-bound elevation | Administrator |
| Labels & terminology | Every user-visible string, per locale and audience | Administrator |
| Saved views | Team views; personal views by the agent | Manager / Agent |
| Resource Hub | Articles, sections, visibility | Manager |

| Area | What can be changed | Permission |
|---|---|---|
| Reports | Custom reports, dashboards, schedules | Manager |

## 12.2 Requires engineering

| Change | Why |
|---|---|
| A third cadence with different mechanics | The two cadences are parameterised; a structurally different one extends the engine. |
| A new condition or action in the rule vocabulary | Operations composes from the catalogue; extending the catalogue is development. |
| A new assignment strategy | Strategies are named implementations. Their parameters are configuration; a new algorithm is not. |
| A new channel adapter | New provider protocol and normalisation. |
| Changing what handle time is derived from | The formula is code; its coefficients are configuration. |
| A new field data type | Storage and rendering. |

The catalogue of conditions and actions is the boundary of Operations' power, and it must be generous at launch, because every gap in it becomes an engineering request.

# Part 13 — The Admin Workspace

The same application, a different role. Eight screens, all scoped to the Onboarding team.

## 13.1 Screen inventory

| Screen | Answers |
|---|---|
| Operations | What needs a person right now? |
| Performance | Is the process working, and did the last change help? |
| All Leads | Fix one thing, or a hundred. |
| Conversation Monitor | What is actually being said? |
| Automation & Cadence | Change the process without a release. |
| Templates & Macros | Control the spend. |
| Configuration | Ticket type, statuses, team, channels, permissions. |
| Audit & Quality | Who changed what, and how well was it handled? |

## 13.2 Bulk actions

Bulk operations need guardrails, not prohibition — a carrier outage or a market pause affects dozens of leads and someone must be able to act once rather than sixty times.

- Maximum selection size per role.
- A confirmation screen stating exactly what will change, and on how many tickets.
- **Bulk outbound messaging is a separate permission from bulk internal changes.**
- Every bulk run is one audit entry carrying the full affected-ticket list, and status and tag changes are reversible within a short window.
- Per-ticket permissions are evaluated individually, and partial failure is reported honestly rather than reported as success for the batch.

## 13.3 Conversation monitoring

Every onboarding conversation, live, filtered by what needs attention. The essential filters are *merchant waiting* (the last message is theirs), *held in queue*, and *template-only*.

From a monitored conversation an administrator may take the ticket over with a reason, add an internal note, or block the contact. They may not reply as themselves (9.2). Blocking stores inbound without creating tickets, sends nothing outbound, leaves existing tickets readable, requires a reason, and is reversible — with both actions audited.

## 13.4 Change management

Five steps, and the live engine never evaluates a draft.

| Step | Behaviour |
|---|---|
| **1 Draft** | Edits accumulate as a pending change. Nothing in production sees them. The interface shows what is unpublished and what the live value still is. |
| **2 Simulate** | Replay against the current queue and the last 30 days, and state the consequence *in business terms* — how many open leads would move to Cold Lead, what the projected template spend becomes, how many walkthrough slots an agent day gains or loses. |
| **3 Publish** | Creates an immutable version with author, timestamp, a diff, and a **full snapshot of the resulting configuration**. |
| **4 Observe** | The rule-health panel flags newly dormant or over-broad rules within 24 hours. The published change is dated on the Performance trend. |
| **5 Roll back** | One click restores a previous version atomically, from its snapshot. |

> **A VERSION IS A SNAPSHOT, NOT A PATCH**
>
> Storing only the diff makes "roll back to v14" restore just the keys v14 happened to touch, silently leaving every later change in place. Store the full resulting configuration with each version. The storage cost is trivial and the alternative is a rollback that does not roll back.

## 13.5 Audit

Configuration changes sit in the *same* immutable log as operational ones, so "why did this lead go cold" and "who changed the cadence" are the same query. Every entry carries actor, timestamp, before and after values, and — where the action required one — the written reason.

# Part 14 — Delivery Plan & Open Decisions

## 14.1 Dependencies on the platform build

This module cannot ship ahead of the platform capabilities it uses. The parent specification places onboarding task types "with their own metrics" in its Phase 3; this document proposes bringing the *module* forward to run alongside Phase 1, because the Zoho pipeline is a live operational cost and the platform pieces it needs are all Phase 1 items.

| Needed from the platform | Platform phase |
|---|---|
| Contacts, conversations, tickets, identity resolution, event log | Phase 1 |
| WhatsApp channel, session window, template registry | Phase 1 |
| Routing engine with the condition and action catalogue | Phase 1 |
| Roles, permissions, field-level rules | Phase 1 |
| Agent workspace: saved views, macros, collision detection, internal notes | Phase 1 |
| Yeastar click-to-call, CDR ingestion, presence | Phase 2 — *the module degrades gracefully without it* |
| Automated QA scoring | Phase 3 — not required |

Without Yeastar the module still works: calls are logged manually by the agent, call duration falls back to the configured default, and the "on a call" capacity signal is unavailable. Nothing else changes.

## 14.2 Phases

| Phase | Scope | Definition of done |
|---|---|---|
| **OA-0**<br>Baseline<br>*1 week* | Record current-state metrics from Zoho before anything changes: registrations per month by source and market, current activation rate, current average time to first contact, current cold rate, current headcount. | Numbers recorded and agreed as the comparison basis. **A project without a baseline cannot be judged.** Nothing blocks this; it can start immediately. |

| Phase | Scope | Definition of done |
|---|---|---|
| **OA-1** Queue & cadence | Ticket type, field set, statuses, capture rules, assignment and capacity, both cadences, agent workspace with the docked conversation, macros, templates, activation and handoff. Zoho cutover. | A WhatsApp registration becomes an assigned onboarding ticket with a running cadence, workable end to end by an agent, activated with a complete handoff record and a full audit trail — **with no engineer involved in the assignment decision**. |
| **OA-2** Admin & measurement | Operations and Performance dashboards, the five metrics, conversation monitoring and take-over, draft → simulate → publish → roll back, audit view, bulk actions. | An administrator changes the cadence from settings, sees the simulated consequence, publishes it, watches the agent view change, and rolls it back — with no deployment. |
| **OA-3** Depth | Handle-time calibration against real CDR data, quality review workspace, Resource Hub feedback loop, win-back evaluation. | Minute values in Appendix D replaced with measured ones; the defaults in this document superseded. |

## 14.3 Open decisions

These are business decisions, not technical ones, and they should be resolved before OA-1 begins. Where a decision is not taken in time, the stated interim default applies and is recorded as technical debt.

| # | Decision | Why it matters | Interim default |
|---|---|---|---|
| 1 | **The monthly WhatsApp template budget for onboarding.** | Five attempts plus three follow-ups is up to eight billable sends per lead. At current volume that is roughly 400 sends a month for onboarding alone. Without a number, the cadence has no design constraint and rule `r5` has nothing to check against. | Ship the ceiling as configuration seeded conservatively, with the per-contact cap enforced, so setting the real number later is a settings change. **Inherits platform decision #3 — resolve them together.** |
| 2 | **The open-lead ceiling per agent.** | Currently the binding constraint on assignment. Set on judgement, not evidence. Too low starves agents; too high recreates the CRM problem of unworkable books. | 12. Review at the end of OA-1 against real handle-time data. |

| # | Decision | Why it matters | Interim default |
|---|---|---|---|
| 3 | **Target activation rate and target time to first contact.** | Without targets, OA-0's baseline has nothing to be compared against and the KPI bands in 11.2 cannot be signed off. | OA-0 measures current values and proposes targets from them; the business confirms before OA-1 exits. |

**WHO CLOSES THESE**

All three are Operations and Commercial decisions. None belongs to the engineering team, and none should be resolved by a developer choosing a default in code. If a decision is still open when the relevant sprint starts, the interim default is implemented *as configuration* so that closing it later is a settings change rather than a rebuild.

# Part 15 — Acceptance Criteria

A feature is not delivered until its criteria pass. "Works on my machine", "the API returns 200" and "the code is merged" are not acceptance. Criteria are **demonstrated, not described**. Two tags are used: CFG means the change must be made from settings by a non-developer with no deployment — if a developer edits a file, the criterion fails even if the behaviour is correct; 2P means two people must be signed in simultaneously.

## Capture and assignment

> **AC-OA-1**
>
> A WhatsApp message to the onboarding number from an unknown number creates a contact, a conversation and one assigned onboarding ticket, with `stage`, `due_date` and `source_utm` populated.
>
> Verify: send from a clean number; confirm one ticket, one conversation, and the assignee's notification.

> **AC-OA-2**
>
> A merchant who messages and then registers through the portal produces exactly **one** onboarding ticket, not two.
>
> Verify: message, then register with the same phone. Confirm a single ticket and a merged contact.

> **AC-OA-3**
>
> Assignment selects the agent with the most free minutes on the lead's market. With two agents differing only in booked walkthroughs, the lead goes to the one with fewer.
>
> Verify: book three walkthroughs for agent A, none for agent B, both on the same market; create a lead; confirm B receives it.

> **AC-OA-4**
>
> **No interface path exists for an agent to select, claim or self-assign a lead**, and the equivalent API call by an Onboarding Agent role is rejected.
>
> Verify: sweep the agent interface for any assignment control; then call the assignment endpoint directly with an agent token.

> **AC-OA-5**
>
> A ticket that cannot be assigned is held with a non-null `hold_reason`, and is assigned automatically within one sweep interval of capacity becoming available. No ticket sits unassigned indefinitely.
>
> Verify: fill every agent to the ceiling, create a lead, confirm the hold and its reason; close one lead; confirm automatic assignment.

## Cadence

**AC-OA-6**

Logging a no-answer outcome increments `attempt_no`, holds the status at Failed to Contact, and schedules the next attempt at the configured gap. At the configured limit the ticket moves to Cold Lead automatically with a logged reason.

Verify: log five consecutive no-answers on one ticket and observe each transition.

**AC-OA-7**

An inbound message resets the cadence exactly **once** per ticket. A second period of silence runs to the cutoff without a further reset.

Verify: reply, go quiet, reply again, go quiet. Confirm one reset and a normal cutoff.

**AC-OA-8**

The went-quiet rule does not send when the template budget is exhausted, and the suppression is visible rather than silent.

Verify: set the ceiling below current usage; confirm no send and a visible reason on the ticket.

**AC-OA-9** CFG

Changing the attempt count from settings changes the cadence for open tickets at the next evaluation, and the change is visible in the agent's cadence tracker without a deployment.

Verify: Operations user at the keyboard, developer forbidden from touching anything.

### Workspace and conversation

**AC-OA-10**

The lead workspace renders the platform conversation component itself. Searching the codebase finds **no** onboarding-specific composer, message renderer or send path.

Verify: code review, plus confirming that a platform change to the composer appears in the onboarding workspace with no onboarding-side change.

**AC-OA-11**

Outside the 24-hour session window, free-text send is blocked **server-side** and approved templates are offered instead. Remaining window time is visible before the agent starts writing.

Verify: force an expired window and attempt a free-text send through the API.

**AC-OA-12**

On a ticket they do not own, an agent can read everything and add an internal note; every other control is disabled **and states its reason**. The equivalent API calls are rejected.

Verify: as an agent, open a colleague's ticket, then call reply, status-change and activate endpoints directly.

**AC-OA-13** 2P

Collision detection warns an agent that a colleague is viewing or replying to the same onboarding conversation, in real time.

Verify: two people, one conversation, both screens visible.

**AC-OA-14**

A macro applies its stated changes and inserts text **without sending**, and states which actions it will take before it runs.

Verify: run each of the six team macros and inspect the resulting ticket and composer.

### Activation and measurement

**AC-OA-15**

Activation is rejected unless all five qualification criteria are true and all eight handoff fields are non-empty — enforced in the service, not only in the interface.

Verify: call the accept endpoint with an incomplete record and confirm rejection.

**AC-OA-16**

Activation transfers contact ownership to the named Account Manager, closes the ticket, and **leaves the conversation open** and attached to the contact.

Verify: activate, then confirm the conversation is intact, visible to the AM, and carries its full pre-activation history.

**AC-OA-17**

Onboarding tickets appear in **no** support volume, SLA, response-time or CSAT aggregate, and receive no satisfaction survey.

Verify: create and solve an onboarding ticket; confirm every support report is unchanged and no survey is sent.

**AC-OA-18**

Every onboarding metric on every screen matches its formula in Appendix C exactly, including which records are excluded, and exports reconcile with the screen.

Verify: recompute each metric from the event log independently and compare.

**AC-OA-19**

The activation-rate trend marks every published configuration change on the time axis with its date and description.

Verify: publish a change, then confirm it appears dated on the chart.

### Administration

**AC-OA-20** CFG

**The whole of Part 12.1 passes.** Work down the table row by row with an Operations user at the keyboard and a developer forbidden from touching anything. Each row is ticked or it is a defect.

**AC-OA-21**

A configuration change can be simulated before it goes live, published as an attributed version, and rolled back atomically — and rolling back to an *earlier* version restores every parameter to that version's values, not only the ones that version changed.

Verify: publish two changes to different parameters, roll back to the version before both, and confirm both are restored.

**AC-OA-22**

An administrator can read any onboarding conversation and take a ticket over with a recorded reason, but has **no reply control** on a ticket they do not own.

Verify: open a colleague's conversation as administrator; confirm no composer; take over and confirm the audit entry names the actor, the previous owner and the reason.

**AC-OA-23**

Bulk actions apply to the selected set only, respect per-ticket permissions individually, report partial failure honestly, and produce one audit entry with the full affected list.

Verify: select a mixed batch including tickets outside the actor's permission; confirm per-item results.

**AC-OA-24**

The full interface works in Arabic with correct right-to-left mirroring and in English, and switching language mid-task loses neither the agent's place, their filters, nor an unsent draft.

Verify: switch language on every screen, mid-task.

## 15.1 Phase exit

| Phase | Exit requires | Blocked by |
|---|---|---|
| OA-0 | Baseline recorded and agreed. | Nothing. Start immediately. |
| OA-1 | AC-OA-1 to 16, and 24. **Without exception, AC-OA-4 and AC-OA-10.** | Decision 1 (template budget) before the cutover date. |
| OA-2 | AC-OA-17 to 23. | Decision 3 (targets) before the KPI bands are signed off. |

**THE TWO CRITERIA THAT MATTER MOST**

If time forces a choice about what to verify carefully, verify **AC-OA-4** and **AC-OA-10**. Everything else in this document describes an onboarding tool. Those two describe *this* one: leads that nobody picks, and a conversation that exists in exactly one place.

# Appendix A — Permission Matrix

Onboarding-relevant grants. Columns are roles as defined in the parent specification, scoped to the Onboarding team. "own" means own tickets; "team" means all onboarding tickets, read.

| Permission | OA | Shift lead | Manager | Admin | Acct mgr | Read-only |
|---|---|---|---|---|---|---|
| View own tickets | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ |
| View all onboarding tickets | **team** | ✓ | ✓ | ✓ | own contacts | ✓ |
| Public reply to contact | **own** | ✓ | ✓ | ✓ | ✓ | — |
| Internal note | ✓ | ✓ | ✓ | ✓ | ✓ | — |
| Edit ticket properties | **own** | ✓ | ✓ | ✓ | ✓ | — |
| Run a macro | **own** | ✓ | ✓ | ✓ | ✓ | — |
| Send approved template | **own** | ✓ | ✓ | ✓ | ✓ | — |
| Reassign / transfer | **—** (request only) | ✓ | ✓ | ✓ | — | — |
| Take over from monitor | — | ✓ | ✓ | ✓ | — | — |
| Assign over capacity | — | ✓ | ✓ | ✓ | — | — |
| Book walkthrough | **own** | ✓ | ✓ | ✓ | — | — |
| Activate seller (accept) | **own**, gated | ✓ | ✓ | ✓ | — | — |
| Reject seller | **—** | — | escalate | ✓ | — | — |
| Reopen cold / junk | **own**, with reason | ✓ | ✓ | ✓ | — | — |
| Bulk internal changes | — | ✓ | ✓ | ✓ | — | — |
| **Bulk outbound messaging** | **—** | — | ✓ | ✓ | — | — |
| Block contact | — | with reason | ✓ | ✓ | — | — |
| Listen to call recordings | own calls | ✓ | ✓ | ✓ | own contacts | — |
| See cost / fee / margin | **—** withheld | — | ✓ | ✓ | ✓ | — |
| Author team macro / template | — | — | ✓ | ✓ | — | — |
| Author personal macro / view | ✓ | ✓ | ✓ | ✓ | ✓ | — |

| Permission | OA | Shift lead | Manager | Admin | Acct mgr | Read-only |
|---|---|---|---|---|---|---|
| Manage automation / cadence | — | — | ✓ | ✓ | — | — |
| Manage roles & permissions | — | — | — | ✓ | — | — |
| View reports | **own** | team | ✓ | ✓ | own contacts | ✓ |
| Export data | — | — | approval | ✓ | — | — |
| View audit log | — | — | ✓ | ✓ | — | ✓ |

*"approval" means the action may be requested and executes once authorised by a role holding that grant. There are no monetary approval caps in this model: authority is granted by role.*

# Appendix B — Rule Catalogue

Ten rules, composed from the platform's existing condition and action catalogue. All are administrator-editable and individually enableable.

### r1 — New registration → Onboarding queue

**WHEN** message.received on Onboarding number · OR · contact.registered

**IF** contact.contact_type = merchant · no open onboarding_task

**THEN** create_ticket(onboarding_task) · set_team(Onboarding) · assign_agent(most_free_minutes, skill: country) · set_field(stage, first_contact) · set_field(due_date, now + first_contact_sla) · notify(assignee)

### r2 — First contact overdue → notify agent and shift lead

**WHEN** schedule.interval · every 15 min

**IF** stage = first_contact · age_business_minutes > first_contact_sla

**THEN** notify(assignee) · notify(team_lead) · add_tag(oa_first_contact_missed)

### r3 — No-answer cadence: daily reminder

**WHEN** schedule.interval · every 15 min

**IF** status = failed_to_contact · attempt_no < no_answer_max · next_attempt_at <= now

**THEN** notify(assignee, "Attempt {{n+1}} of {{max}} due") · set_field(due_date, now)

### r4 — No-answer cutoff → Cold Lead

**WHEN** schedule.interval · every 15 min

**IF** status = failed_to_contact · attempt_no >= no_answer_max · time_since_last_inbound > 24h

**THEN** set_status(cold_lead) · add_tag(oa_cold_no_answer) · add_internal_note(reason) · clear(due_date) · notify(assignee)

### r5 — Went quiet → follow-up template

**WHEN** ticket.status_age_exceeded · awaiting reply > quiet_gap

**IF** followup_no < quiet_max · contact.status != opted_out · template_budget.remaining > 0 · contact.template_sends < per_contact_cap

**THEN** send_message(whatsapp, template: oa_nudge_v2) · set_field(followup_no, +1) · set_field(next_followup_at, now + quiet_gap) · set_field(due_date, next_followup_at)

### r6 — Went-quiet cutoff → Cold Lead

**WHEN** schedule.interval · every 15 min

**IF** followup_no >= quiet_max · status in (prequalified, qualified) · time_since_last_inbound > quiet_gap

**THEN** set_status(cold_lead) · add_tag(oa_cold_went_quiet) · clear(due_date) · notify(assignee)

### r7 — Reply resets the cadence, once

**WHEN** message.received · inbound

**IF** ticket.type = onboarding_task · cadence_reset = false

**THEN** set_field(attempt_no, 0) · set_field(followup_no, 0) · set_field(cadence_reset, true) · set_field(due_date, now) · notify(assignee)

### r8 — Holding queue: assign when capacity frees

**WHEN** schedule.interval · every 15 min · OR · agent.capacity_changed

**IF** assignee is null · status is open · any agent can_take

**THEN** assign_agent(most_free_minutes, skill: country) · clear(hold_reason) · notify(assignee) · notify(team_lead if held > escalation_threshold)

### r9 — Day-30 win-back, then archive · *ships disabled*

**WHEN** schedule.interval · daily 09:00

**IF** status = cold_lead · days_since_status_change >= winback_day · contact.status = active · template_budget.remaining > 0

**THEN** send_message(whatsapp, template: oa_winback_d30) · add_tag(oa_winback_sent)

### r10 — First order placed → close the onboarding ticket

**WHEN** domain_event · order.created

**IF** contact has onboarding_task with status = activated · first_order_at is null

**THEN** set_field(first_order_at, event.time) · solve_ticket() · notify(account_manager)

# Appendix C — Metric Formula Reference

Normative computation reference. Every screen must match these exactly, including exclusions. Durations are business hours unless stated.

```
activation_rate
    = count(status = activated)
    ÷ count(status in (activated, cold_lead))
      where  created_at >= window_start
        and  age_at_terminal <= 14 days
        and  status != junk_lead                      ← junk excluded from BOTH sides

speed_to_first_contact
    = count(first_outbound_at − created_at <= first_contact_sla)
    ÷ count(tickets created in window)
      business hours only; computed from the event log,
      NOT from platform first_response_time

clean_handoff_rate
    = count(activations where all 8 handoff fields non-empty)
    ÷ count(activations)

first_order_7d                                        ← computed in MDM reporting
    = count(first_order_at − activated_at <= 7 days)
    ÷ count(activations)

funnel_hygiene
    = count(open tickets where due_date < now)
    ÷ count(open tickets)                             sampled at end of day; lower is
better

cold_rate
    = count(status = cold_lead) ÷ count(tickets, excl. junk_lead)

handle_time            see 10.3
free_minutes           see 4.2
cohort_activation(d)
    = count(activated within d days of created_at)
    ÷ count(registrants in that week, excl. junk_lead)
```

**Reporting rules.** Report the median and the 90th percentile for any duration, never the mean. Publish the denominator beside every rate. Compare an agent only to their own history and to the onboarding-team median.

# Appendix D — Configuration Parameter Reference

Every parameter, its launch default and its owner. Defaults marked *estimate* are to be replaced with measured values at OA-3.

| Parameter | Default | Owner | Notes |
|---|---|---|---|
| **Cadence** | | | |
| `first_contact_sla` | 24 h | Manager | Business hours |
| `no_answer_max` | 5 | Manager | Attempts before Cold Lead |
| `no_answer_gap` | 1 day | Manager | |
| `quiet_max` | 3 | Manager | Each follow-up is billable |
| `quiet_gap` | 3 days | Manager | |
| `stalled_day` | 7 | Manager | Written reason required |
| `winback_day` | 30 | Manager | r9 disabled at launch |
| **Capacity** | | | |
| `shift_minutes` | 480 | Manager | 08:30–16:30 |
| `reserved_minutes` | 60 | Manager | Morning triage + end-of-day close |
| `max_open_leads` | 12 | Manager | **Open decision 2** — currently the binding constraint |
| `holding_sweep_interval` | 15 min | `manage_automation` | |
| `holding_escalation` | 60 min | `manage_automation` | Notify shift lead beyond this |
| **Time model — all *estimate*** | | | |
| `walkthrough_minutes` | 20 | Manager | The call itself |
| `walkthrough_buffer` | 5 | Manager | Notes and handoff record |
| `first_contact_minutes` | 6 | Manager | |
| `wrap_minutes` | 3 | Manager | After any answered call |
| `no_answer_minutes` | 1.5 | Manager | Dial, wait, log |
| `followup_minutes` | 4 | Manager | One cadence touch with its note |

| Parameter | Default | Owner | Notes |
|---|---|---|---|
| `compose_minutes` | 2 | Manager | Free-text outbound |
| `read_minutes` | 0.5 | Manager | Inbound message |
| `template_minutes` | 0.5 | Manager | Pre-written |
| `note_minutes` | 1.5 | Manager | |
| `activation_minutes` | 12 | Manager | Accept action + handoff record |
| **Templates** | | | |
| `template_budget_month` | seeded low | Manager | **Open decision 1** |
| `template_per_contact_cap` | 8 | Manager | Hard stop regardless of cadence |
| `marketing_templates` | prohibited | Administrator | On the onboarding number |
| **Assignment** | | | |
| `assignment_strategy` | most_free_minutes | `manage_automation` | |
| `assignment_skill` | contact.country | `manage_automation` | Soft — falls back to the whole team |
| `assignment_tiebreak` | fewest_open_leads | `manage_automation` | |

# Appendix E — Decision Log

| # | Decision | Status | Consequence |
|---|---|---|---|
| 1 | Onboarding is a ticket type inside the support platform, not a separate CRM or application. | **Settled** | Determines the whole of Parts 1, 2 and 6. Reversing it means rebuilding, not reconfiguring. |
| 2 | Zoho CRM is decommissioned rather than integrated. | **Settled** | Single cutover, no parallel run (3.5). |
| 3 | The platform assigns; agents never pick. | **Settled** | No claim control anywhere. AC-OA-4. |
| 4 | Assignment is by free minutes, not round-robin or lead count. | **Settled** | Requires the time model (Part 10) to exist before assignment can work. |
| 5 | Repeat contact attempts are a counter on one status, not separate statuses. | **Settled** | Seven statuses, and a cutoff that is one numeric condition. |
| 6 | The conversation is docked into the lead workspace, not linked to. | **Settled** | No second composer. AC-OA-10. |
| 7 | The walkthrough is 20 minutes, not 45. | **Settled** | Roughly doubles walkthrough slots per agent day. Process detail relocates to the Resource Hub. |
| 8 | Handle time is derived from the ticket, never measured by a timer or self-reported. | **Settled** | Consistent, un-gameable, directional. Must not enter individual performance scoring (10.4). |
| 9 | Onboarding carries `counts_toward: none` and no SLA. | **Settled** | Inherited from the platform specification. AC-OA-17. |
| 10 | Day-30 win-back ships built but disabled. | **Settled** | Enabling it is a configuration change, to be made against cohort evidence. |
| 11 | Monthly WhatsApp template budget for onboarding. | **Open** | Caps the cadence. Interim: seeded low with the per-contact cap enforced. |
| 12 | Open-lead ceiling per agent. | **Open** | Currently the binding constraint on assignment. Interim: 12. |
| 13 | Target activation rate and target time to first contact. | **Open** | Blocks sign-off of the KPI bands. Interim: derive from the OA-0 baseline. |

**END OF SPECIFICATION**

Questions, disagreements and defects against this document go to the Commercial Director for process questions and to the platform owner for anything that touches the parent specification. A discrepancy between this document and the Omnichannel Support Platform specification is always a defect here, never there.